One real query: count sports items among 2,048 news records

GLM | prefix cache ON | source count = 517

1 Check labels

80 single-row reference calls

+ 80 batch tests (eight rows each)

80 / 80 comparisons match

Choose batch size 8

Checking: 160 requests

2 Execute every row

256 new batches cover all 2,048 rows

Returned COUNT = 519

17 / 2,048 labels differ from S1/S2

0.83% reference loss < 5% tolerance

Deployment: 256 requests

3 Add the whole bill

Requests: 160 + 256 = 416

S1 uses 2,048 requests

Tokens: 93,120 + 255,425

= 348,545 vs 338,365 for S1

Calls: 79.69% lower

Raw tokens: 3.01% higher

In this execution, the accepted batch meets observed fidelity tolerance; checking adds to its cost.